

Transfer learning for high-dimensional linear regression

Notes

19331053 纪传宇

2022 年 7 月 15 日

1 INTRODUCTION

(1)Regression analysis is one of the most widely used statistical methods to understand the association of an outcome with a set of covariates In this paper, we consider transfer learning in high-dimensional linear models. Formally, the target model can be written as

$$y_i^{(0)} = (x_i^{(0)})^T \beta + \epsilon_i^{(0)}, i = 1, \dots, n_0 \quad (1)$$

where $((x_i^{(0)})^T, y_i^{(0)}), i = 1, \dots, n_0$ are independent samples, $\beta \in R^p$ is the coefficient vector of interest, and $\epsilon_i^{(0)}, i = 1, \dots, n_0$ are independently distributed random noises with $E[\epsilon_i^{(0)} | x_i^{(0)}] = 0$. In the high-dimensional regime, where p can be larger and much larger than n_0 , β is often assumed to be sparse such that the number of nonzero elements of β , denoted by s , is much smaller than p .

(2) In the context of transfer learning, we observe additional samples from K auxiliary studies, That is, we observe $((x_i^{(k)})^T, y_i^{(k)})$ generated from the auxiliary model

$$y_i^{(k)} = (x_i^{(k)})^T w^{(k)} + \epsilon_i^{(k)}, i = 1, \dots, n_k, k = 1, \dots, K \quad (2)$$

where $w^{(k)} \in R^p$ is the regression vector for the k th study, and $\epsilon_i^{(k)}$ is the random noise such that $E[\epsilon_i^{(k)} | x_i^{(k)}] = 0$. The regression coefficients $w^{(k)}$ are unknown and different from our target β in general. The number of auxiliary studies, K , is allowed to grow but practically K may not be too large. We will study the estimation and prediction of target model (1) utilizing the primary data $y_i^{(0)} = (x_i^{(0)})^T \beta + \epsilon_i^{(0)}, i = 1, \dots, n_0$ as well as the data from K auxiliary studies $y_i^{(k)} = (x_i^{(k)})^T w^{(k)} + \epsilon_i^{(k)}, i = 1, \dots, n_k, k = 1, \dots, K$

If an auxiliary model is ‘similar’ to the target model, we say that this auxiliary sample/study is informative. In this work, we characterize the informative level of the k th auxiliary study using the sparsity of the difference between $w^{(k)}$ and β . Let $\delta^{(k)} = \beta - w^{(k)}$ denote the contrast between $w^{(k)}$ and β . The set of informative auxiliary samples is those whose contrasts are sufficiently sparse:

$$\mathcal{A}_q = \left\{ 1 \leq k \leq K : \|\delta^{(k)}\|_q \leq h \right\} \quad (3)$$

for some $q \in [0, 1]$. The set \mathcal{A}_q contains the auxiliary studies whose contrast vectors have ℓ_q -sparsity at most h and is called the informative set. It will be seen later that as long as h is relatively small compared to the sparsity of β , the studies in \mathcal{A}_q can be useful in improving the prediction and estimation of β . For any $q \in [0, 1]$, smaller h implies that the auxiliary samples in \mathcal{A}_q are more informative; larger cardinality of \mathcal{A}_q ($|\mathcal{A}_q|$) implies that a larger number of informative auxiliary samples. Therefore, smaller h and larger $|\mathcal{A}_q|$ should be favourable. We allow \mathcal{A}_q to be empty in which case none of the auxiliary samples is informative. For the auxiliary samples outside of \mathcal{A}_q , we do not assume sparse $\delta^{(k)}$ and hence $w^{(k)}$ can be very different from β for $k \notin \mathcal{A}_q$.

2 Notation

2.1 title

表 1: Notation for Trans-Lasso algorithm

<u><i>Indices</i></u>	
$X^{(0)} \in R^{n_0 \times p}$	design matrix for the primary data
$Y^{(0)} \in R^{n_0}$	response vector for the primary data
$X^{(k)} \in R^{n_k \times p}$	design matrix for the kth auxiliary data
$Y^{(k)} \in R^{n_k}$	the response vector for the kth auxiliary data
$\{R_l\}_{l \in \mathcal{L}}$	$R_l, l \in \mathcal{L}$
$n_{\mathcal{A}_q}$	$\sum_{k \in \mathcal{A}_q} n_k$
$\Lambda_{\max}(\Sigma)$	largest eigenvalues of Σ respectively
$\Lambda_{\min}(\Sigma)$	smallest eigenvalues of Σ respectively
$a \vee b$	$\max\{a, b\}$
$a \wedge b$	$\min\{a, b\}$
$a_n = O(b_n)$ and $a_n \lesssim b_n$	$ a_n/b_n \leq c$ for some constant c when n is large enough
$a_n \asymp b_n$	$ a_n/b_n \rightarrow c$ for some constant c as $n \rightarrow \infty$
$a_n = O_P(b_n)$ and $a_n \lesssim_{\mathbb{P}} b_n$	$\mathbb{P}(a_n/b_n \leq c) \rightarrow 1$ for some constant $c < \infty$
$a_n = o_P(b_n)$	$(a_n/b_n > c) \rightarrow 0$ for any constant $c > 0$

3 ESTIMATION WITH KNOWN INFORMATIVE AUXILIARY SAMPLES

3.1 Oracle Trans-Lasso algorithm

Algorithm 1: Oracle Trans-Lasso algorithm

Input : Primary data $(X^{(0)}, y^{(0)})$ and informative auxiliary samples $\{X^{(k)}, y^{(k)}\}_{k \in \mathcal{A}}$

Output: $\hat{\beta}$

Step 1. Compute

$$\hat{w}^{\mathcal{A}} = \arg \min_{w \in \mathbb{R}^p} \left\{ \frac{1}{2n_{\mathcal{A}}} \sum_{k \in \mathcal{A}} \|y^{(k)} - X^{(k)}w\|_2^2 + \lambda_w \|w\|_1 \right\} \quad (4)$$

for $\lambda_w = c_1 \sqrt{\log p / n_{\mathcal{A}}}$ with some constant c_1 .

Step 2. Let

$$\hat{\beta} = \hat{w}^{\mathcal{A}} + \hat{\delta}^{\mathcal{A}}, \quad (5)$$

where

$$\hat{\delta}^{\mathcal{A}} = \arg \min_{\delta \in \mathbb{R}^p} \left\{ \frac{1}{2n_0} \|y^{(0)} - X^{(0)}(\hat{w}^{\mathcal{A}} + \delta)\|_2^2 + \lambda_{\delta} \|\delta\|_1 \right\} \quad (6)$$

for $\lambda_{\delta} = c_2 \sqrt{\log p / n_0}$ with some constant c_2 .

Fig. 1. algorithm of Oracle Trans-Lasso algorithm

3.2 step of Oracle Trans-Lasso algorithm

\mathcal{A} is known in this algorithm

(1) compute an initial estimator using all the informative auxiliary samples $\hat{w}^{\mathcal{A}}$ is realized based on the Lasso (Tibshirani, 1996) using all the informative auxiliary samples

to find $\hat{w}^{\mathcal{A}}$ in

$$\hat{w}^{\mathcal{A}} = \arg \min_{w \in \mathbb{R}^p} \left\{ \frac{1}{2n_{\mathcal{A}}} \sum_{k \in \mathcal{A}} \|y^{(k)} - X^{(k)}w\|_2^2 + \lambda_w \|w\|_1 \right\} \quad (4)$$

is equivalent to find

$$\hat{w}^{\mathcal{A}} = \arg \min_{w \in \mathbb{R}^p} \left\{ \frac{1}{2n_{\mathcal{A}}} \sum_{k \in \mathcal{A}} w^T x^{(k)T} x^{(k)} w - 2y^{(k)T} x^{(k)} w + y^{(k)T} y \right\} \quad (5)$$

let $f(w)$ denote

$$\left\{ \frac{1}{2n_{\mathcal{A}}} \sum_{k \in \mathcal{A}} w^T x^{(k)T} x^{(k)} w - 2y^{(k)T} x^{(k)} w + y^{(k)T} y \right\} \quad (6)$$

then we can use maximum likelihood estimates to evaluate $\hat{w}^{\mathcal{A}}$ if we plug $\hat{w}^{\mathcal{A}}$ into the formula we can get $E[\frac{\partial f(w)}{\partial w} |_{\hat{w}^{\mathcal{A}}}] = 0$ which is equal to

$$E \left[\frac{1}{n_{\mathcal{A}}} \sum x^{(k)\top} (x^{(k)} w - y^{(k)}) \right] = 0 \quad (7)$$

so the equation in the article

$$\mathbb{E} \left[\sum_{k \in \mathcal{A}} (X^{(k)})^T (y^{(k)} - X^{(k)} w^{\mathcal{A}}) \right] = 0 \quad (8)$$

can be proved

Denoting $\mathbb{E} \left[x_i^{(k)} (x_i^{(k)})^T \right] = \Sigma^{(k)}$, $w^{\mathcal{A}}$ has the following explicit form:

$$w^{\mathcal{A}} = \beta + \delta^{\mathcal{A}} \quad (9)$$

the probabilistic limit of $\widehat{W}^{\mathcal{A}}$, has bias $\delta^{\mathcal{A}}$, and $\delta^{\mathcal{A}}$ can be calculated by

$$\delta^{\mathcal{A}} = \sum_{k \in \mathcal{A}} n_k / n_{\mathcal{A}} \times \delta^{(k)} \quad (10)$$

(2) Correct its bias using the primary data in the second step $\delta^{\mathcal{A}}$ is a sparse high-dimensional vector while ℓ_1 -norm is no larger than h .

Because we set $\mathcal{A}_q = \{1 \leq k \leq K : \|\delta^{(k)}\|_q \leq h\}$ and in this algorithm q is 1. What's more, $\|\delta^{(k)}\|_1 \leq h$ and $\delta^{\mathcal{A}} = \sum_{k \in \mathcal{A}} n_k / n_{\mathcal{A}} \times \delta^{(k)}$.

Hence, the error of step 2 is under control for a relatively small h

The Oracle Trans-Lasso does not penalize the differences among the regression coefficients in the auxiliary studies. This is again because the focus of transfer learning is only the target study. Theoretically, extra penalization terms and the joint analysis of multiple estimators may not help improve the estimation accuracy of the parameter of interest.

4 UNKNOWN SET OF INFORMATIVE AUXILIARY SAMPLES

we propose a data-driven method for estimation and prediction when \mathcal{A} is unknown.

4.1 The Trans-Lasso algorithm

two main step: (1) construct a collection of candidate estimators, each of which is based on an estimate of \mathcal{A} .

(2) perform an aggregation step on these candidate estimators

more notation: For a generic estimate of \mathcal{A} , b , denote its sum of squared prediction error as

$$\widehat{Q}(\mathcal{I}, b) = \sum_{i \in \mathcal{I}} \left\| y_i^{(0)} - \left(x_i^{(0)} \right)^\top b \right\|_2^2 \quad (11)$$

\mathcal{I} is a subset of $\{1, \dots, n_0\}$. Let $\Lambda^{L+1} = \left\{ v \in \mathbb{R}^{L+1} : v_l \geq 0, \sum_{l=0}^L v_l = 1 \right\}$ denote an L-dimensional simplex

Algorithm 2: Trans-Lasso Algorithm

Input : Primary data $(X^{(0)}, y^{(0)})$ and samples from K auxiliary studies $\{X^{(k)}, y^{(k)}\}_{k=1}^K$.

Output: $\hat{\beta}^{\hat{\theta}}$.

Step 1. Let \mathcal{I} be a random subset of $\{1, \dots, n_0\}$ such that $|\mathcal{I}| \approx c_0 n_0$ with some constant $0 < c_0 < 1$. Let $\mathcal{I}^c = \{1, \dots, n_0\} \setminus \mathcal{I}$.

Step 2. Construct $L + 1$ candidate sets of \mathcal{A} , $\{\hat{G}_0, \hat{G}_1, \dots, \hat{G}_L\}$ such that $\hat{G}_0 = \emptyset$ and $\hat{G}_1, \dots, \hat{G}_L$ are based on (14) using $(X_{\mathcal{I}^c}^{(0)}, y_{\mathcal{I}^c}^{(0)})$ and $\{X^{(k)}, y^{(k)}\}_{k=1}^K$.

Step 3. For each $0 \leq l \leq L$, run the Oracle Trans-Lasso algorithm with primary sample $(X_{\mathcal{I}^c}^{(0)}, y_{\mathcal{I}^c}^{(0)})$ and auxiliary samples $\{X^{(k)}, y^{(k)}\}_{k \in \hat{G}_l}$. Denote the output as $\hat{\beta}(\hat{G}_l)$ for $0 \leq l \leq L$.

Step 4. Compute

$$\hat{\theta} = \tag{10}$$

$$\arg \min_{\theta \in \Lambda^{L+1}} \left\{ \hat{Q}(\mathcal{I}^c, \sum_{l=0}^L \hat{\beta}(\hat{G}_l) \theta_l) + \sum_{l=0}^L \theta_l \hat{Q}(\mathcal{I}^c, \hat{\beta}(\hat{G}_l)) + \frac{2\lambda_\theta}{n_0} \sum_{l=0}^L \theta_l \log(\theta_l) \right\}$$

for some $\lambda_\theta > 0$. Output

$$\hat{\beta}^{\hat{\theta}} = \sum_{l=0}^L \hat{\theta}_l \hat{\beta}(\hat{G}_l). \tag{11}$$

Fig. 2. algorithm of Trans-Lasso algorithm

steps 2 and 3 of the Trans-Lasso algorithm construct some initial estimates of $\beta, \hat{\beta}(\hat{G}_l)$. They are computed using the Oracle Trans-Lasso algorithm by treating each \hat{G}_l as the set of informative auxiliary samples.

Step 4 is based on the Q-aggregation proposed in Dai et al. (2012) with a uniform prior, a Kullback-Leibler penalty, and a simplified tuning parameter. The Q-aggregation can be viewed as a weighted version of least square aggregation and exponential aggregation (Rigollet Tsybakov, 2011) and it has been shown to be rate optimal both in expectation and with high probability for model selection aggregation problems

4.2 Constructing the candidate sets for aggregation

Model selection aggregation is an effective method for the transfer learning task under consideration

let us first point out a naive construction of candidate sets, which consists of 2^K candidates. These candidates are all different combinations of $1, \dots, K$ denoted by $\widehat{G}_1, \dots, \widehat{G}_{2^K}$. \mathcal{A} is an element of these candidate sets. However, the number of candidates is too large and it can be computationally burdensome. In contrast, we would like to pursue a much smaller number of candidate sets such that the cost of aggregation is almost negligible and

$$\mathbb{P}\left(\widehat{G}_l \subseteq \mathcal{A}, \text{ for some } 1 \leq l \leq L\right) \rightarrow 1 \quad (12)$$

can be achieved under mild conditions.

the ideas to solve this problem is to exploit the sparsity patterns of the contrast vectors. The definition of \mathcal{A} implies that $\{\delta^{(k)}\}_{k \in \mathcal{A}}$ are sparser than $\{\delta^{(k)}\}_{k \in \mathcal{A}^c}$, where $\mathcal{A}^c = \{1, \dots, K\} \setminus \mathcal{A}$. This property motivates us to find a sparsity index $R^{(k)}$ and its estimator $\widehat{R}^{(k)}$ for each $1 \leq k \leq K$ such that

$$\max_{k \in \mathcal{A}^0} R^{(k)} < \min_{k \in \mathcal{A}^c} R^{(k)} \quad \text{and} \quad \mathbb{P}\left(\max_{k \in \mathcal{A}^0} \widehat{R}^{(k)} < \min_{k \in \mathcal{A}^c} \widehat{R}^{(k)}\right) \rightarrow 1, \quad (13)$$

where \mathcal{A}^0 is some subset of \mathcal{A} . In words, the sparsity indices in \mathcal{A}^0 are no larger than the sparsity indices in \mathcal{A}^c and so are their estimators with high probability. To utilize Equation (13), we can define the candidate sets as

$$\widehat{G}_l = \left\{ 1 \leq k \leq K : \widehat{R}^{(k)} \text{ is among the first } l \text{ smallest of all } \right\} \quad (14)$$

for $1 \leq l \leq K$. That is, \widehat{G}_l is the set of auxiliary samples whose estimated sparsity indices are among the first l smallest. A direct consequence of Equations (13) and (14) is that $\mathbb{P}\left(\widehat{G}_{|\mathcal{A}^0|} = \mathcal{A}^0\right) \rightarrow 1$ and

To achieve the largest gain with transfer learning, we would like to find proper sparsity indices such that Equation (13) holds for $\sum_{k \in \mathcal{A}^0} n_k$ as large as possible. Notice that $\widehat{G}_{K+1} = \{1, \dots, K\}$ is always included as candidates according to Equation (14). Hence, in the special cases where all the auxiliary samples are informative or none of the auxiliary samples are informative, it

holds that $\widehat{G}_{|\mathcal{A}|} = \mathcal{A}$ and the Trans-Lasso is not much worse than the Oracle Trans-Lasso. The more challenging cases are $0 < |\mathcal{A}| < K$.

As $\{\delta^{(k)}\}_{k \in \mathcal{A}^c}$ are not necessarily sparse, the estimation of $\delta^{(k)}$ or functions of $\delta^{(k)}$, $1 \leq k \leq K$, is not trivial. As an example, an intuitive sparsity index can be $\|\delta^{(k)}\|_1$ and its estimate is $\|\widehat{\beta}(\widehat{G}_0) - \widehat{w}^{(k)}\|_1$, where $\widehat{w}^{(k)}$ is the Lasso estimate of $w^{(k)}$ based on the k th study. However, such a Lasso-based estimate is not guaranteed to converge to the oracle $\|\delta^{(k)}\|_1$ when $\delta^{(k)}$ is non-sparse. Therefore, we consider using $R^{(k)} = \|\Sigma\delta^{(k)}\|_2^2$, which is a function of the population-level marginal statistics, as the oracle sparsity index for k th auxiliary sample. The advantage of $R^{(k)}$ is that it has a natural unbiased estimate even when $\delta^{(k)}$ is non-sparse. Let us relate $R^{(k)}$ to the sparsity of $\delta^{(k)}$ using a Bayesian characterization of sparse vectors assuming $\Sigma^{(k)} = \Sigma$ for all $0 \leq k \leq K$. If $\delta_j^{(k)}$ are i.i.d. Laplacian distributed with mean zero and variance v_k^2 for each k , then it follows from the properties of Laplacian distribution (Liu & Kozubowski, 2015) that $\mathbb{E}[\|\delta^{(k)}\|_1] \asymp \mathbb{E}^{1/2}[\|\Sigma\delta^{(k)}\|_2^2]$. Hence, the rank of $\mathbb{E}[\|\Sigma\delta^{(k)}\|_2^2]$ is the same as the rank of $\mathbb{E}[\|\delta^{(k)}\|_1]$. As $\max_{k \in \mathcal{A}} \|\delta^{(k)}\|_1 < \min_{k \in \mathcal{A}^c} \|\delta^{(k)}\|_1$, it is reasonable to expect $\max_{k \in \mathcal{A}} \|\Sigma\delta^{(k)}\|_2^2 < \min_{k \in \mathcal{A}^c} \|\Sigma\delta^{(k)}\|_2^2$. The above derivation holds for many other zero mean prior distributions besides Laplacian. This illustrates our motivation for considering $R^{(k)}$ as the oracle sparsity index.

We next introduce the estimated version $\widehat{R}^{(k)}$, based on the primary data $\left\{ \left(x_i^{(0)} \right)^\top, y_i^{(0)} \right\}_{i \in \mathcal{I}}$ (after sample splitting) and auxiliary samples $\{X^{(k)}, y^{(k)}\}_{k=1}^K$. We first perform a SURE screening (Fan & Lv, 2008) on the marginal statistics to reduce the effects of random noises.

In this article, if we have the linear regression:

$$y_i = \beta_0 + x_{1i}^\top \beta_1 + \cdots + x_{pi}^\top \beta_p + \varepsilon_i, \quad i = 1, \dots, n \quad (15)$$

when the number of variables p is much greater than n , how to do with this data. Fan and Lv consider that the marginal correlation coefficient of x_{ij} and y_j that $\text{corr}(y_i, x_{ij}) := \omega_j$ should be highly correlated with the regression coefficient β_j .

So if β_j is small, the marginal correlation coefficient of x_{ij} and y_j should be very small in most cases. So before using Lasso in variable selection for

an ultra-high dimensional regression model, we should use some measures related to marginal correlation coefficient to rank the marginal correlation coefficients. And only the covariates corresponding to the top $t_* = n_*^\alpha$ correlation coefficients (absolute values of) are selected and put into the regression model

We summarize our proposal for Step 2 of the Trans-Lasso as follows (Algorithm 3). Let $n_* = \min_{0 \leq k \leq K} n_k$.

One can see that $\widehat{\Delta}^{(k)}$ are empirical marginal statistics such that $\mathbb{E} [\widehat{\Delta}^{(k)}] = \Sigma \delta^{(k)}$ for $k \in \mathcal{A}$. The set \widehat{T}_k is the set of first t_* largest marginal statistics for the k th sample. The purpose of screening the marginal statistics is to reduce the magnitude of noise. Notice that the un-screened version $\|\widehat{\Delta}^{(k)}\|_2^2$ is a sum of p random variables and it contains noise of order $p / (n_k \wedge n_0)$, which diverges fast as p is much larger than the sample sizes. By screening with t_* of order $n_*^{\alpha, \alpha < 1}$, the errors induced by the random noises is under control. In practice, the auxiliary samples with very small sample sizes can be removed from the analysis as their contributions to the target problem is mild. Desirable choices of \widehat{T}_k should keep the variation of $\Sigma \delta^{(k)}$ as much as possible. Under proper conditions, SURE screening can consistently select a set of strong marginal statistics and hence is appropriate for the current purpose. In Step 2.2, we compute $\widehat{R}^{(k)}$ based on the marginal statistics which are selected by SURE screening. In practice, different choices of t_* may lead to different realizations of \widehat{G}_l . One can compute multiple sets of $\{\widehat{R}^{(k)}\}_{k=1}^K$ with different t_* which give multiple sets of $\{\widehat{G}_l\}_{l=1}^K$. It will be seen from Lemma 1 that a finite number of choices on t_* does not affect the rate of convergence.

Algorithm 3: Step 2 of the Trans-Lasso Algorithm

Step 2.1. For $1 \leq k \leq K$, compute the marginal statistics

$$\hat{\Delta}^{(k)} = \frac{1}{n_k} \sum_{i=1}^{n_k} x_i^{(k)} y_i^{(k)} - \frac{1}{|\mathcal{I}|} \sum_{i \in \mathcal{I}} x_i^{(0)} y_i^{(0)}. \quad (15)$$

For each $k \in \{1, \dots, K\}$, let \hat{T}_k be obtained by SURE screening such that

$$\hat{T}_k = \left\{ 1 \leq j \leq p : |\hat{\Delta}_j^{(k)}| \text{ is among the first } t_* \text{ largest of all} \right\}$$

for a fixed $t_* = n_*^\alpha$, $0 \leq \alpha < 1$.

Step 2.2. Define the estimated sparse index for the k -th auxiliary sample as

$$\hat{R}^{(k)} = \left\| \hat{\Delta}_{\hat{T}_k}^{(k)} \right\|_2^2. \quad (16)$$

Step 2.3. Compute \hat{G}_l as in (14) for $l = 1, \dots, L$.

Fig. 3. step 2 of the Trans-Lasso algorithm

4.3 Theoretical properties of Trans-Lasso

we first establish the model selection aggregation type of results for the Trans-Lasso estimator

Lemma 1 (Q-aggregation for Trans-Lasso). Assume that Conditions 1 and 2 hold true. Let $\hat{\theta}$ be computed via Equation (10) with $\lambda_\theta \geq 4\sigma_0^2$. With probability at least $1 - t$, it holds that

$$\frac{1}{|\mathcal{I}^c|} \left\| X_{\mathcal{I}^c, \cdot}^{(0)} (\hat{\beta} - \beta) \right\|_2^2 \leq \min_{0 \leq l \leq L} \frac{1}{|\mathcal{I}^c|} \left\| X_{\mathcal{I}^c, \cdot}^{(0)} \left(\hat{\beta} \left(\hat{G}_l \right) - \beta \right) \right\|_2^2 + \frac{\lambda_\theta \log(L/t)}{n_0}. \quad (16)$$

If $L \leq c_1 n_0$ for some small enough constant c_1 , then

$$\left\| \hat{\beta}^{\hat{\theta}} - \beta \right\|_2^2 \lesssim_{\mathbb{P}} \min_{0 \leq l \leq L} \left\| \hat{\beta} \left(\hat{G}_l \right) - \beta \right\|_2^2 + \frac{\log L}{n_0} \quad (17)$$

Lemma 1 implies that the performance of $\hat{\beta}^{\hat{\theta}}$ only depends on the best candidate regardless of the performance of other candidates under mild conditions. As the

original Lasso is always in our dictionary, Equations (16) and (17) imply that $\hat{\beta}^{\hat{\theta}}$ is not much worse than the Lasso in prediction and estimation. Formally, ‘not much worse’ refers to the last term in Equation (16), which can be viewed as the cost of ‘searching’ for the best candidate model within the dictionary which is of order $\log L/n_0$. This term is almost negligible, say, when $L = O(K)$, which corresponds to our constructed candidate estimators. This demonstrates the robustness of $\hat{\beta}^{\hat{\theta}}$ to adversarial auxiliary samples.

Indeed, an estimator with ℓ_2 -error guarantee is crucial for more challenging tasks, such as out-of-sample prediction and inference. For our transfer learning task, we show in Equation (17) that the estimation error is of the same order if the cardinality of the dictionary is $L \leq cn_0$ for some small enough c . For our constructed dictionary, it suffices to require $K \leq cn_0$. In many practical applications, K is relatively small compared to the sample sizes and hence this assumption is not very restrictive.

For each $k \in \mathcal{A}^c$, define a set

$$H_k = \left\{ 1 \leq j \leq p : \left| \Sigma_{j,\cdot}^{(k)} w^{(k)} - \Sigma_{j,\cdot}^{(0)} \beta \right| > n_*^{-\kappa}, \kappa < \alpha/2 \right\}. \quad (18)$$

Recall that $\alpha < 1$ is defined such that $t_* = n^\alpha$. In fact, H_k is the set of ‘strong’ marginal statistics that can be consistently selected into \hat{T}_k for each $k \in \mathcal{A}^c$. We see that $\Sigma_{j,\cdot}^{(k)} w^{(k)} - \Sigma_{j,\cdot}^{(0)} \beta = \Sigma_{j,\cdot} \delta^{(k)}$ if $\Sigma^{(k)} = \Sigma^{(0)}$ for $k \in \mathcal{A}^c$. The definition of \mathcal{H}_k in Equation (18) allows for heterogeneous designs among non-informative auxiliary samples.

5 SIMULATION STUDIES

we evaluate the empirical performance of the proposed methods and some other comparable methods in various numerical experiments. Specifically, we evaluate the performance of five methods, including Lasso, Oracle Trans-Lasso proposed in Section 2.1, TransLasso proposed in Section 3.1, and two other ad hoc transfer learning methods related to ours

- (1)'Lasso'
- (2)'Oracle Trans-Lasso'
- (3)'Trans-Lasso'
- (4)'aggregated Lasso', which is Trans-Lasso except the bias-correction step (Step 2) of the Oracle Trans-Lasso
- (5)'ad hoc ℓ_1 -transfer', which follows the steps of Trans-Lasso but uses a different aggregation step. Considering $\widehat{R}^{(k)} = \left\| \widehat{\beta}^L - \widehat{w}^{(k)} \right\|_1, k = 1, \dots, K$, where $\widehat{\beta}^L$ and $\widehat{w}^{(k)}$ are the Lasso estimators based on each of the corresponding studies. Moreover, the Q-aggregation step is replaced with the cross-validation, where we select the set \widehat{G}_l that minimizes the out-of-sample prediction errors.

5.1 Retrieval results with identity covariance matrix for the designs

We consider $p=500, n_0 = 150$, and $n_1, \dots, n_K = 100$ for $K=20$. The covariates $x_i^{(k)}$ are i.i.d. Gaussian with mean zero and identity covariance matrix for all $0 \leq k \leq K$ and $\epsilon_i^{(k)}$ are i.i.d. Gaussian with mean zero and variance one for all $0 \leq k \leq K$. For the target parameter β , we set $s=16$, $\beta_j = 0.3$ for $j \in \{1, \dots, s\}$, and $\beta_j = 0$ otherwise. For the regression coefficients in auxiliary samples, we consider two configurations.

- (i) For a given \mathcal{A} , if $k \in \mathcal{A}$, let

$$w_j^{(k)} = \beta_j - 0.31 \quad (j \in H_k) \tag{19}$$

where H_k is a random subset of $[p]$ with $|H_k| = h \in \{2, 6, 12\}$. If $k \notin \mathcal{A}$, we set H_k to be a random subset of $[p]$ with $|H_k| = 2s$ and $w_j^{(k)} = \beta_j - 0.51 \quad (j \in H_k)$. We set $w_1^{(k)} = -0.3$ for $k = 1, \dots, K$.




Fig. 4. Estimation errors of the ad hoc ℓ_1 -transfer, Agg-Lasso, Lasso, Oracle Trans-Lasso, and Trans-Lasso with identity covariance matrices of the predictor in (i)




Fig. 5. Estimation errors of the ad hoc ℓ_1 -transfer, Agg-Lasso, Lasso, Oracle Trans-Lasso, and Trans-Lasso with identity covariance matrices of the predictor in (ii)

(ii) For a given \mathcal{A} , if $k \in \mathcal{A}$, let $H_k = \{1, \dots, 100\}$ and

$$w_j^{(k)} = \beta_j + \xi_j \mathbb{1}(k \in H_k), \quad \text{where } \xi_j \sim_{i.i.d.} N(0, h/100), \quad (20)$$

where $h \in \{2, 6, 12\}$ and $N(a, b)$ is the normal with mean a and standard deviation b . If $k \notin \mathcal{A}$, we set $H_k = \{1, \dots, 100\}$ and

$$w_j^{(k)} = \beta_j + \xi_j \mathbb{1}(j \in H_k), \quad \text{where } \xi_j \sim_{i.i.d.} N(0, 2s/100) \quad (21)$$

We set $w_1^{(k)} = -0.3$ for $k = 1, \dots, K$. The setting (i) can be treated as either ℓ_0 - or ℓ_1 -sparse contrasts. In practice, the true parameters are unknown and we use \mathcal{A} to denote the set of auxiliary samples without distinguishing ℓ_0 - or ℓ_1 -sparsity. We consider $|\mathcal{A}| \in \{0, 4, 8, \dots, 20\}$

In the performance, we can clearly see that the Lasso does not change as $|\mathcal{A}|$ increases. The other algorithms based on transfer learning have estimation errors decreasing as $|\mathcal{A}|$ increases. In settings (i) and (ii), the Oracle Trans-Lasso has the smallest estimation errors in most settings. In the article the proposed Trans-Lasso is always the second best but in my code repetition, Agg-Lasso almost as same as Trans-Lasso. When $\mathcal{A} = \emptyset$, the Trans-Lasso can have smaller errors than the oracle Trans-Lasso where the latter one does not use auxiliary information. This implies that some auxiliary information can still be borrowed. Due to the randomness of the parameter generation, our definition of \mathcal{A} may not always be the best subset of auxiliary samples that give the smallest estimation errors.

The article mention that Agg-Lasso method has larger estimation errors than Trans-Lasso and ad hoc ℓ_1 -transfer. But in my repetition, Agg-Lasso's SSE is not too different from Trans-Lasso.

5.2 Retrieval results with homogeneous designs among $\mathcal{A} \cup \{0\}$

We now consider $x_i^{(k)}$ as i.i.d. Gaussian with mean zero and a equi-correlated covariance matrix, where $\Sigma_{j,j} = 1$ and $\Sigma_{j,k} = 0.8$ if $j \neq k$ for $k \in \mathcal{A} \cup \{0\}$. For $k \notin \mathcal{A} \cup \{0\}$, $x_i^{(k)}$ are i.i.d. Gaussian with mean zero and a Toeplitz covariance matrix whose first row is

$$\Sigma_{1,.}^{(k)} = (1, \underbrace{1/(k+1), \dots, 1/(k+1)}_{2k-1}, 0_{p-2k}) \quad (22)$$

Other true parameters and the dimensions of the samples are set to be the same as in Section 5.1.




Fig. 6. Estimation errors of the ad hoc ℓ_1 -transfer, Agg-Lasso, Lasso, Oracle Trans-Lasso, and Trans-Lasso with identity covariance matrices of the predictor in (i)




Fig. 7. Estimation errors of the ad hoc ℓ_1 -transfer, Agg-Lasso, Lasso, Oracle Trans-Lasso, and Trans-Lasso with identity covariance matrices of the predictor in (ii)

From the results presented in Fig 6 and Fig 7, except (ii) configurations and $|h| = 12$, the Trans-Lasso and Oracle Trans-Lasso have reliable performance in the current setting. However the SSE are large in this type then in Section 5.1 because of the highly correlated covariates. When h is relatively large, according to the article Agg-Lasso and ad hoc ℓ_1 -transfer have significantly larger estimation errors than Trans-Lasso. However, during repetition, we can see that Agg-Lasso's SSE still is not too different from Trans-Lasso. But the figure also demonstrates the advantage of Trans-Lasso over some ad hoc methods

5.3 Retrieval results with heterogeneous designs

We next consider a setting where $\Sigma^{(k)}$ are distinct for $k = 0, \dots, K$. Specifically, for $k = 1, \dots, K$, let $x_i^{(k)}$ as i.i.d. Gaussian with mean zero and a Toeplitz covariance matrix whose first row is Equation (22). Moreover, $\Sigma^{(0)} = I_p$. Other parameters and the dimensions of the samples are set to be the same as in Section 5.1




Fig. 8. Estimation errors of the ad hoc ℓ_1 -transfer, Agg-Lasso, Lasso, Oracle Trans-Lasso, and Trans-Lasso with identity covariance matrices of the predictor in (i)




Fig. 9. Estimation errors of the ad hoc ℓ_1 -transfer, Agg-Lasso, Lasso, Oracle Trans-Lasso, and Trans-Lasso with identity covariance matrices of the predictor in (ii)

From the results presented in Fig8 and Fig9,the Trans-Lasso and Oracle Trans-Lasso and the Agg-Lasso have reliable performance in the current

setting. In some situation, Agg-Lasso's SSE is smaller than Trans-Lasso which is different from the article's demonstration

5.4 SIMULATION STUDIES CODE

The R code for all the methods

```

1
2  library(glmnet)
3  agg.fun<- function(B, X.test,y.test , total.step=10,
4  selection=F){
5      if(sum(B==0)==ncol(B)*nrow(B)){
6          return(rep(0,nrow(B)))
7      }
8      p<-nrow(B)
9      K<-ncol(B)
10     colnames(B)<-NULL
11     if(selection){#select beta.hat with smallest prediction
12         error
13         khat<-which.min(colSums((y.test-X.test%*%B)^2))
14         theta.hat<-rep(0, ncol(B))
15         theta.hat[khat] <- 1
16         beta=B[,khat]
17         beta.ew=NULL
18     }else{#Q-aggregation
19         theta.hat<- exp(-colSums((y.test-X.test%*%B)^2)/2)
20         theta.hat=theta.hat/sum(theta.hat)
21         theta.old=theta.hat
22         beta<-as.numeric(B%*%theta.hat)
23         beta.ew<-beta
24         # theta.old=theta.hat
25         for(ss in 1:total.step){
26             theta.hat<- exp(-colSums((y.test-X.test%*%B)^2)/2+
27             colSums((as.vector(X.test%*%beta)-X.test%*%B)^2)/8)
28             theta.hat<-theta.hat/sum(theta.hat)
29             beta<- as.numeric(B%*%theta.hat*1/4+3/4*beta)
30             if(sum(abs(theta.hat-theta.old))<10^(-3)){break}
31             theta.old=theta.hat

```

```

29     }
30   }
31   list(theta=theta.hat, beta=beta, beta.ew=beta.ew)
32 }
33
34
35 ###oracle Trans-Lasso
36 las.kA<-function(X, y, A0, n.vec, lam.const=NULL, l1=T){
37   p<-ncol(X)
38   size.A0<- length(A0)
39   if(size.A0 > 0){
40     ind.kA<- ind.set(n.vec, c(1, A0+1))
41     ind.l<-1:n.vec[1]
42     if(l1){
43       y.A<-y[ind.kA]
44     }else{ #the l0-method
45       y.A<- y[ind.l]
46       Sig.hat<-t(X)%*%X/nrow(X)
47       for(k in 1:size.A0){
48         ind.k<- ind.set(n.vec,k+1)
49         lam.k <- sqrt(mean(y[ind.l]^2)/n.vec[1]+mean(y[ind.k
50 ]^2)/n.vec[k]) * sqrt(2*log(p))
51         delta.hat.k<-lassoshooting(XtX=Sig.hat,
52           Xty=t(X[ind.k,])%*%y[ind.k]/n.vec[k+1]-t(X[1:n.vec
53 [1],])%*%y[1:n.vec[1]]/n.vec[1],
54           lambda=lam.k)$coef
55         y.A<-c(y.A, y[ind.k]-X[ind.k,]%*%delta.hat.k)
56       }
57     }
58     if(is.null(lam.const)){
59       cv.init<-cv.glmnet(X[ind.kA,], y.A, nfolds=8, lambda=
60 seq(1,0.1,length.out=10)*sqrt(2*log(p)/length(ind.kA)))
61       lam.const <- cv.init$lambda.min/sqrt(2*log(p)/length(

```

```

ind.kA)))
62      # cv.delta<-cv.glmnet(x=X[ind.1,],y=y[ind.1]-X[ind.1,]%*
%w.kA, lambda=seq(1,0.1,length.out=10)*sqrt(2*log(p)/length(
ind.1)))
63      #delta.kA<-predict(cv.delta, s='lambda.min', type='
coefficients')[−1]
64      delta.kA <- as.numeric(glmnet(x=X[ind.1,],y=y[ind.1]-X[
ind.1,]%*w.kA, lambda=lam.const*sqrt(2*log(p)/length(ind.1)))
$beta)
65      delta.kA<-delta.kA*(abs(delta.kA)>=lam.const*sqrt(2*log(
p)/length(ind.1)))
66      beta.kA <- w.kA + delta.kA
67      lam.const=NA
68    }else{
69      cv.init<-cv.glmnet(X[1:n.vec[1],], y[1:n.vec[1]], nfolds
=8, lambda=seq(1,0.1,length.out=20)*sqrt(2*log(p)/n.vec[1]))
70      lam.const<-cv.init$lambda.min/sqrt(2*log(p)/n.vec[1])
71      beta.kA <- predict(cv.init, s='lambda.min', type='
coefficients')[−1]
72      w.kA<-NA
73    }
74    list(beta.kA=as.numeric(beta.kA),w.kA=w.kA, lam.const=lam.
const)
75
76  }
77
78  #Trans Lasso method
79  Trans.lasso <- function(X, y, n.vec, I.til, ll=T){
80    M= length(n.vec)−1
81    #step 1
82    X0.til<-X[I.til,] #used for aggregation
83    y0.til<-y[I.til]
84    X<- X[−I.til,]
85    y<-y[−I.til]
86    #step 2
87    Rhat <- rep(0, M+1)
88    p<- ncol(X)
89    n.vec[1]<- n.vec[1]−length(I.til)

```

```

90     ind.1<-ind.set(n.vec,1)
91     for(k in 2:(M+1)){
92         ind.k<-ind.set(n.vec,k)
93         Xty.k <- t(X[ind.k,])%*%y[ind.k]/n.vec[k] - t(X[ind.1,])
94         %*%y[ind.1]/n.vec[1]
95         margin.T<-sort(abs(Xty.k),decreasing=T)[1:round(n.vec[1]
96         /3)]
97         Rhat[k] <- sum(margin.T^2)
98     }
99     Tset<-list()
100     k0=0
101     kk.list<-unique(rank(Rhat[-1]))
102     #cat(rank(Rhat[-1]),'\n')
103     for(kk in 1:length(kk.list)){#use Rhat as the selection
104     rule
105         Tset[[k0+kk]]<-which(rank(Rhat[-1]) <= kk.list[kk])
106     }
107     k0=length(Tset)
108     Tset<-unique(Tset)
109     #cat(length(Tset),'\n')
110
111     beta.T<-list()
112     init.re<-las.kA(X=X, y=y, A0=NULL, n.vec=n.vec, l1=l1)
113     beta.T[[1]] <- init.re$beta.kA
114     beta.pool.T<-beta.T ##another method for comparison
115     for(kk in 1:length(Tset)){#use pi.hat as selection rule
116         T.k <- Tset[[kk]]
117         re.k<- las.kA(X=X, y=y, A0=T.k, n.vec=n.vec, l1=l1, lam.
118         const=init.re$lam.const)
119         beta.T[[kk+1]] <-re.k$beta.kA
120         beta.pool.T[[kk+1]]<-re.k$w.kA
121     }
122     beta.T<-beta.T[!duplicated((beta.T))]
123     beta.T<- as.matrix(as.data.frame(beta.T))
124     agg.re1 <- agg.fun(B=beta.T, X.test=X0.til, y.test=y0.til)
125     beta.pool.T<-beta.pool.T[!duplicated((beta.pool.T))]
126     beta.pool.T<- as.matrix(as.data.frame(beta.pool.T))
127     agg.re2<-agg.fun(B=beta.pool.T, X.test=X0.til, y.test=y0.

```

```

124 til)
125     return(list(beta.hat=agg.re1$beta, theta.hat=agg.re1$theta
, rank.pi=rank(Rhat[-1]),
126     beta.pool=agg.re2$beta, theta.pool=agg.re2$theta))
127 }
128
129
130 #A method for comparison: Trans-Lasso(l1). It has the same
pipeline of the Trans-Lasso
131 ###but with sparsity index  $R_k = \|\hat{w}^{(k)} - \beta\|_1$  and a
naive aggregation (empirical risk minimization)
132 Trans.lasso.sp <- function(X, y, n.vec, I.til, l1=T){
133   M= length(n.vec)-1
134   #step 1
135   X0.til<-X[I.til,] #used for aggregation
136   y0.til<-y[I.til]
137   X<- X[-I.til,]
138   y<-y[-I.til]
139   #step 2
140   Rhat <- rep(0, M+1)
141   p<- ncol(X)
142   n.vec[1]<- n.vec[1]-length(I.til)
143   ind.l<-ind.set(n.vec,1)
144   init.re<-las.kA(X=X, y=y, A0=NULL, n.vec=n.vec, l1=l1)
145   for(k in 2: (M+1)){
146     ind.k <- ind.set(n.vec,k)
147     w.init.k<- as.numeric(glmnet(X[ind.k,], y[ind.k], lambda
=init.re$lam.const*sqrt(2*log(p)/length(ind.k)))$beta)
148     Rhat[k] <- sum(abs(w.init.k-init.re$beta.kA)) ## $\|\hat{w}^{(k)} - \beta\|_1$ 
149   }
150   Tset<- list()
151   k0=0
152   kk.list<-unique(rank(Rhat[-1]))
153   #cat(rank(Rhat[-1]),'\n')
154   for(kk in 1:length(kk.list)){#use pi.hat as selection rule
155     Tset[[k0+kk]]<- which(rank(Rhat[-1]) <= kk.list[kk])

```

```

156     }
157     k0=length(Tset)
158     Tset<- unique(Tset)
159     # cat(length(Tset),'\n')
160
161     beta.T<-list()
162     beta.T[[1]] <- init.re$beta.kA
163     for(kk in 1:length(Tset)){#use pi.hat as selection rule
164         T.k <- Tset[[kk]]
165         beta.T[[kk+1]] <- las.kA(X=X, y=y, A0=T.k, lam.const=init
.re$lam.const, n.vec=n.vec, ll=ll)$beta.kA
166     }
167     beta.T<-beta.T[!duplicated((beta.T))]
168     beta.T<- as.matrix(as.data.frame(beta.T))
169     agg.re <- agg.fun(B=beta.T, X.test=X0.til, y.test=y0.til,
selection =T)
170     return(list(beta.sp=agg.re$beta, theta.sp=agg.re$theta,
rank.pi=rank(Rhat[-1])))
171 }
172
173 #computing the MSE
174 mse.fun<- function(beta,est, X.test=NULL){
175     pred.err<-NA
176     est.err<- sum((beta-est)^2)
177
178     if(!is.null(X.test)){
179         pred.err<- mean((X.test%*%(beta-est))^2)
180     }
181     return(list(est.err=est.err, pred.err= pred.err))
182 }
183
184 ind.set<- function(n.vec, k.vec){
185     ind.re <- NULL
186     for(k in k.vec){
187         if(k==1){
188             ind.re<-c(ind.re,1: n.vec[1])
189         }else{
190             ind.re<- c(ind.re, (sum(n.vec[1:(k-1)])+1): sum(n.vec

```

```

    [1:k]))
191     }
192   }
193   ind.re
194 }
195 rep.col<-function(x,n){
196   matrix(rep(x,each=n), ncol=n, byrow=TRUE)
197 }
198
199

```

Code for plot

```

1   library(tidyverse)
2
3
4   Coef.gen<- function(s, h,q=30, size.A0, M, sig.beta,sig.
delta1, sig.delta2, p, exact=T){
5     beta0 <- c(rep(sig.beta,s), rep(0, p - s))
6     W<- rep.col(beta0, M)# ten prior estimates
7     W[1,]<-W[1,]-2*sig.beta
8     for(k in 1:M){
9       if(k <= size.A0){
10        if(exact){
11          samp0<- sample(1:p, h, replace=F)
12          W[samp0,k] <-W[samp0,k] + rep(-sig.delta1, h)
13        }else{
14          W[1:100,k] <-W[1:100,k] + rnorm(100, 0, h/100)
15        }
16      }else{
17        if(exact){
18          samp1 <- sample(1:p, q, replace = F)
19          W[samp1,k] <- W[samp1,k] + rep(-sig.delta2,q)
20        }else{
21          W[1:100,k] <-W[1:100,k] + rnorm(100, 0, q/100)
22        }
23      }
24    }
25  }

```



```

24     }
25     return(list(W=W, beta0=beta0))
26 }
27
28
29 fig1 <- function(h,size.A0, exact = F){
30     p = 500
31     s = 16
32     M = 20
33     sig.beta = 0.3
34
35     sig.z <- 1
36     n0 <- 150
37     M = 20
38     n.vec <- c(n0, rep(100, M))
39     Sig.X <- diag(1, p)
40     Niter = 500
41     ll=T
42
43     A0 = 1:size.A0
44     beta0<-
45     coef.all <-Coef.gen( s, h = h, q = 2*s, size.A0 = size.A0,
46     M = M, sig.beta = sig.beta,
47     sig.delta1 = sig.beta, sig.delta2 = sig.beta+0.2, p = p,
48     exact=exact)
49     B <- cbind(coef.all$beta0, coef.all$W)
50     beta0 <- coef.all$beta0
51     library(mvtnorm)
52     ###generate the data
53     X <- NULL
54     y <- NULL
55     for (k in 1:(M + 1)) {
56         X <- rbind(X, rmvnorm(n.vec[k], rep(0, p), Sig.X))
57         ind.k <- ind.set(n.vec, k)
58         y <- c(y, X[ind.k, ] %*% B[, k] + rnorm (n.vec[k], 0, 1)
59     )
60     }
61     ###compute init beta ####

```

```

59     mse.vec<-rep(NA,6)
60     beta.init <-
61     as.numeric(glmnet(X[1:n.vec[1], ], y[1:n.vec[1]], lambda =
sqrt(2 * log(p) / n.vec[1]))$beta)
62     mse.vec[1] = mse.fun(as.numeric(beta.init), beta0)$est.err
63
64     #####Oracle Trans-Lasso#####
65     if (size.A0 == 0) {
66         beta.kA <- beta.init
67     } else{
68         beta.kA <- las.kA(X, y, A0 = 1:size.A0, n.vec = n.vec,
l1=l1)$beta.kA
69     }
70     mse.vec[2] = mse.fun(as.numeric(beta.kA), beta0)$est.err
71     #####Trans-Lasso#####
72     prop.re1 <- Trans.lasso(X, y, n.vec, I.til = 1:50, l1 = l1
)
73     prop.re2 <- Trans.lasso(X, y, n.vec, I.til = 101:n.vec[1],
l1=l1)
74     if(size.A0 > 0 & size.A0< M){ ##Rank.re characterizes the
performance of the sparsity index Rk
75         Rank.re<- (sum(prop.re1$rank.pi[1:size.A0]<=size.A0) +
76         sum(prop.re2$rank.pi[1:size.A0]<=size.A0))/2/size.A0
77     }else{ Rank.re <- 1 }
78     beta.prop <- (prop.re1$beta.hat + prop.re2$beta.hat) / 2
79     mse.vec[3] = mse.fun(beta.prop, beta0)$est.err
80
81     #####A method for comparison: it has the same pipeline of
the Trans-Lasso
82     #####but with sparsity index  $R_k = \|\hat{w}^{(k)} - \beta\|_1$  and a
naive aggregation (empirical risk minimization)
83     prop.sp.re1 <- Trans.lasso.sp(X, y, n.vec, I.til = 1:50,
l1 = l1)
84     prop.sp.re2 <- Trans.lasso.sp(X, y, n.vec, I.til = 101:n.
vec[1], l1=l1)
85     if(size.A0 > 0 & size.A0< M){
86         Rank.re.sp <- (sum(prop.sp.re1$rank.pi[1:size.A0]<=size.
A0) +

```

```

87     sum(prop.sp.re2$rank.pi[1:size.A0]<=size.A0))/2/size.A0
88   }else{ Rank.re.sp <-1 }
89   beta.sp <- (prop.sp.re1$beta.sp + prop.sp.re2$beta.sp) / 2
90   mse.vec[4] = mse.fun(beta.sp, beta0)$est.err
91
92   #####another method for comparison: it is the same as
Trans-Lasso except
93   ##that the bias correction step (step 2 of Oracle Trans-
Lasso) is omitted
94   beta.pool<-(prop.re1$beta.pool+prop.re2$beta.pool)/2
95   mse.vec[5] = mse.fun(beta.pool, beta0)$est.err
96   #####naive translasso: simply assumes A0=1K
97   beta.all <- las.kA(X, y, A0 = 1:M, n.vec = n.vec, l1=l1)$
beta.kA#naive transLasso
98   mse.vec[6] = mse.fun(as.numeric(beta.all), beta0)$est.err
99   return(mse.vec)
100 }
101
102 h = c(2,6,12)
103 size.A0 = c(0,4,8,12,16,20)
104
105 plot.sse = function(data,i){
106   data %>% gather(key = "method", value = "SSE", -A) %>%
107   ggplot(aes(x = A, y = SSE, group = method, color = method,
shape = method)) +
108   geom_point()+
109   geom_line() +
110   ylab(paste0("SSE (h=",i,")")) +
111   theme_bw()
112 }
113
114 plot.data <- function(sse,i){
115   data = sse[[i]][,-6]
116   data = as.data.frame(data)
117   names(data) = c("Lasso", "Oracle-Trans-Lasso", "Trans-Lasso",
,"Agg-Lasso", "hc-l1")
118   data$A = size.A0
119   plot.sse(data, h[i])

```

```

120 }
121
122 sse11 = list()
123
124 for (i in 1:3){
125   ssevec = matrix(NA, ncol = 6, nrow = 6)
126   for(j in 1:6){
127     ssevec[j, ] = fig1(h[i], size.A0[j])
128   }
129   sse11[[i]] = ssevec
130 }
131
132 library(ggpubr)
133
134 p11.1 = plot.data(sse11, 1)
135 p11.2 = plot.data(sse11, 2)
136 p11.3 = plot.data(sse11, 3)
137
138 ggarrange(p11.1, p11.2, p11.3, ncol=3, nrow=1, common.legend
= TRUE, legend="bottom")
139
140 sse12 = list()
141 for (i in 1:3){
142   ssevec = matrix(NA, ncol = 6, nrow = 6)
143   for(j in 1:6){
144     ssevec[j, ] = fig1(h[i], size.A0[j], exact = T)
145   }
146   sse12[[i]] = ssevec
147 }
148
149 p12.1 = plot.data(sse12, 1)
150 p12.2 = plot.data(sse12, 2)
151 p12.3 = plot.data(sse12, 3)
152
153 ggarrange(p12.1, p12.2, p12.3, ncol=3, nrow=1, common.legend
= TRUE, legend="bottom")
154
155 fig2 <- function(h, size.A0, exact = F){

```

```

156     p = 500
157     s = 16
158     M = 20
159     sig.beta = 0.3
160
161     sig.z <- 1
162     n0 <- 150
163     M = 20
164     n.vec <- c(n0, rep(100, M))
165     Sig.X <- diag(1, p)
166     Niter = 200
167     ll=T
168
169     A0 = 1:size.A0
170     beta0<-
171     coef.all <-Coef.gen( s, h = h, q = 2*s, size.A0 = size.A0,
172     M = M, sig.beta = sig.beta,
173     sig.delta1 = sig.beta, sig.delta2 = sig.beta+0.2, p = p,
174     exact=exact)
175     B <- cbind(coef.all$beta0, coef.all$W)
176     beta0 <- coef.all$beta0
177     library(mvtnorm)
178     ###generate the data
179     X <- NULL
180     y <- NULL
181     for (k in 1:(M + 1)) {
182         if(k<=size.A0+1)
183         {Sig.X[which(Sig.X == 0)] = 0.8}
184         else{
185             Sig.X[1, (1+1):(2*k-1+1)] = 1/(k+1)
186             Sig.X[(1+1):(2*k-1+1),1] = 1/(k+1)}
187         X <- rbind(X, rmvnorm(n.vec[k], rep(0, p), Sig.X))
188         ind.k <- ind.set(n.vec, k)
189         y <- c(y, X[ind.k, ] %*% B[, k] + rnorm(n.vec[k], 0, 1)
190     )
191     }
192     ###compute init beta ####
193     mse.vec<-rep(NA,6)

```

```

191     beta.init <-
192     as.numeric(glmnet(X[1:n.vec[1], ], y[1:n.vec[1]], lambda =
sqrt(2 * log(p) / n.vec[1]))$beta)
193     mse.vec[1] = mse.fun(as.numeric(beta.init), beta0)$est.err
194
195     #####Oracle Trans-Lasso#####
196     if (size.A0 == 0) {
197         beta.kA <- beta.init
198     } else{
199         beta.kA <- las.kA(X, y, A0 = 1:size.A0, n.vec = n.vec,
l1=l1)$beta.kA
200     }
201     mse.vec[2] = mse.fun(as.numeric(beta.kA), beta0)$est.err
202     #####Trans-Lasso#####
203     prop.re1 <- Trans.lasso(X, y, n.vec, I.til = 1:50, l1 = l1
)
204     prop.re2 <- Trans.lasso(X, y, n.vec, I.til = 101:n.vec[1],
l1=l1)
205     if(size.A0 > 0 & size.A0< M){ #Rank.re characterizes the
performance of the sparsity index Rk
206         Rank.re<- (sum(prop.re1$rank.pi[1:size.A0]<=size.A0) +
207         sum(prop.re2$rank.pi[1:size.A0]<=size.A0))/2/size.A0
208     }else{ Rank.re <- 1 }
209     beta.prop <- (prop.re1$beta.hat + prop.re2$beta.hat) / 2
210     mse.vec[3] = mse.fun(beta.prop, beta0)$est.err
211
212     #####A method for comparison: it has the same pipeline of
the Trans-Lasso
213     #####but with sparsity index  $R_k = \|\hat{w}^{(k)} - \beta\|_1$  and a
naive aggregation (empirical risk minimization)
214     prop.sp.re1 <- Trans.lasso.sp(X, y, n.vec, I.til = 1:50,
l1 = l1)
215     prop.sp.re2 <- Trans.lasso.sp(X, y, n.vec, I.til = 101:n.
vec[1], l1=l1)
216     if(size.A0 > 0 & size.A0< M){
217         Rank.re.sp <- (sum(prop.sp.re1$rank.pi[1:size.A0]<=size.
A0) +
218         sum(prop.sp.re2$rank.pi[1:size.A0]<=size.A0))/2/size.A0

```

```

219     }else{ Rank.re.sp <-1 }
220     beta.sp <- (prop.sp.re1$beta.sp + prop.sp.re2$beta.sp) / 2
221     mse.vec[4] = mse.fun(beta.sp, beta0)$est.err
222
223     #####another method for comparison: it is the same as
Trans-Lasso except
224     ##that the bias correction step (step 2 of Oracle Trans-
Lasso) is omitted
225     beta.pool<-(prop.re1$beta.pool+prop.re2$beta.pool)/2
226     mse.vec[5] = mse.fun(beta.pool, beta0)$est.err
227     #####naive translasso: simply assumes A0=1:K
228     beta.all <- las.kA(X, y, A0 = 1:M, n.vec = n.vec, l1=l1)$
beta.kA##naive transLasso
229     mse.vec[6] = mse.fun(as.numeric(beta.all), beta0)$est.err
230     return(mse.vec)
231   }
232
233   sse21 = list()
234
235   for (i in 1:3){
236     ssevec = matrix(NA, ncol = 6, nrow = 6)
237     for(j in 1:6){
238       ssevec[j, ] = fig2(h[i], size.A0[j])
239     }
240     sse21[[i]] = ssevec
241   }
242
243   p21.1 = plot.data(sse21, 1)
244   p21.2 = plot.data(sse21, 2)
245   p21.3 = plot.data(sse21, 3)
246   ggarrange(p21.1, p21.2, p21.3, ncol=3, nrow=1, common.legend
= TRUE, legend="bottom")
247
248   sse22 = list()
249   for (i in 1:3){
250     ssevec = matrix(NA, ncol = 6, nrow = 6)
251     for(j in 1:6){
252       ssevec[j, ] = fig2(h[i], size.A0[j], exact = T)

```

```

253     }
254     sse22[[i]] = ssevec
255 }
256
257 p22.1 = plot.data(sse22, 1)
258 p22.2 = plot.data(sse22, 2)
259 p22.3 = plot.data(sse22, 3)
260 ggarrange(p22.1, p22.2, p22.3, ncol=3, nrow=1, common.legend
= TRUE, legend="bottom")
261
262 fig3 <- function(h, size.A0, exact = F){
263     p = 500
264     s = 16
265     M = 20
266     sig.beta = 0.3
267
268     sig.z <- 1
269     n0 <- 150
270     M = 20
271     n.vec <- c(n0, rep(100, M))
272     Sig.X <- diag(1, p)
273     Niter = 200
274     ll=T
275
276     A0 = 1:size.A0
277     beta0<-
278     coef.all <- Coef.gen( s, h = h, q = 2*s, size.A0 = size.A0,
M = M, sig.beta = sig.beta,
279     sig.delta1 = sig.beta, sig.delta2 = sig.beta+0.2, p = p,
exact=exact)
280     B <- cbind(coef.all$beta0, coef.all$W)
281     beta0 <- coef.all$beta0
282     library(mvtnorm)
283     ###generate the data
284     X <- NULL
285     y <- NULL
286     X <- rbind(X, rmvnorm(n.vec[1], rep(0, p), Sig.X))
287     ind.1 <- ind.set(n.vec, 1)

```



```

288 y <- c(y, X[ind.1, ] %*% B[, 1] + rnorm (n.vec[1], 0, 1))
289 for (k in 2:(M + 1)) {
290   Sig.X[1, (1+1):(2*k-1+1)] = 1/(k+1)
291   Sig.X[(1+1):(2*k-1+1), 1] = 1/(k+1)
292   X <- rbind(X, rmvnorm(n.vec[k], rep(0, p), Sig.X))
293   ind.k <- ind.set(n.vec, k)
294   y <- c(y, X[ind.k, ] %*% B[, k] + rnorm (n.vec[k], 0, 1)
295 )
296   }
297   ###compute init beta####
298   mse.vec<-rep(NA,6)
299   beta.init <-
300   as.numeric(glmnet(X[1:n.vec[1], ], y[1:n.vec[1]], lambda =
301   sqrt(2 * log(p) / n.vec[1]))$beta)
302   mse.vec[1] = mse.fun(as.numeric(beta.init), beta0)$est.err
303   #####Oracle Trans-Lasso#####
304   if (size.A0 == 0) {
305     beta.kA <- beta.init
306   } else{
307     beta.kA <- las.kA(X, y, A0 = 1:size.A0, n.vec = n.vec,
308     l1=l1)$beta.kA
309   }
310   mse.vec[2] = mse.fun(as.numeric(beta.kA), beta0)$est.err
311   #####Trans-Lasso#####
312   prop.re1 <- Trans.lasso(X, y, n.vec, I.til = 1:50, l1 = l1
313 )
314   prop.re2 <- Trans.lasso(X, y, n.vec, I.til = 101:n.vec[1],
315   l1=l1)
316   if(size.A0 > 0 & size.A0<M){ ##Rank.re characterizes the
317   performance of the sparsity index Rk
318   Rank.re<- (sum(prop.re1$rank.pi[1:size.A0]<=size.A0) +
319   sum(prop.re2$rank.pi[1:size.A0]<=size.A0))/2/size.A0
320   }else{ Rank.re <- 1 }
321   beta.prop <- (prop.re1$beta.hat + prop.re2$beta.hat) / 2
322   mse.vec[3] = mse.fun(beta.prop, beta0)$est.err
323
324   #####A method for comparison: it has the same pipeline of

```

```

320   the Trans-Lasso
      ###but with sparsity index  $R_k = \|\hat{w}^{(k)} - \beta\|_1$  and a
naive aggregation (empirical risk minimization)
321   prop.sp.re1 <- Trans.lasso.sp(X, y, n.vec, I.til = 1:50,
l1 = l1)
322   prop.sp.re2 <- Trans.lasso.sp(X, y, n.vec, I.til = 101:n.
vec[1], l1=l1)
323   if(size.A0 > 0 & size.A0< M){
324     Rank.re.sp <- (sum(prop.sp.re1$rank.pi[1:size.A0]<=size.
A0) +
325       sum(prop.sp.re2$rank.pi[1:size.A0]<=size.A0))/2/size.A0
326   }else{ Rank.re.sp <-1 }
327   beta.sp <- (prop.sp.re1$beta.sp + prop.sp.re2$beta.sp) / 2
328   mse.vec[4] = mse.fun(beta.sp, beta0)$est.err
329
330   #####another method for comparison: it is the same as
Trans-Lasso except
331   ##that the bias correction step (step 2 of Oracle Trans-
Lasso) is omitted
332   beta.pool<-(prop.re1$beta.pool+prop.re2$beta.pool)/2
333   mse.vec[5] = mse.fun(beta.pool, beta0)$est.err
334   #####naive translasso: simply assumes A0=1:K
335   beta.all <- las.kA(X, y, A0 = 1:M, n.vec = n.vec, l1=l1)$
beta.kA#naive transLasso
336   mse.vec[6] = mse.fun(as.numeric(beta.all), beta0)$est.err
337   return(mse.vec)
338 }
339
340 sse31 = list()
341
342 for (i in 1:3){
343   ssevec = matrix(NA, ncol = 6, nrow = 6)
344   for(j in 1:6){
345     ssevec[j, ] = fig3(h[i], size.A0[j])
346   }
347   sse31[[i]] = ssevec
348 }
349

```

```

350     p31.1 = plot.data(sse31, 1)
351     p31.2 = plot.data(sse31, 2)
352     p31.3 = plot.data(sse31, 3)
353     ggarrange(p31.1, p31.2, p31.3, ncol=3, nrow=1, common.legend
= TRUE, legend="bottom")
354
355     sse32 = list()
356     for (i in 1:3){
357         ssevec = matrix(NA, ncol = 6, nrow = 6)
358         for(j in 1:6){
359             ssevec[j, ] = fig3(h[i], size.A0[j], exact = T)
360         }
361         sse32[[i]] = ssevec
362     }
363
364     p32.1 = plot.data(sse32, 1)
365     p32.2 = plot.data(sse32, 2)
366     p32.3 = plot.data(sse32, 3)
367     ggarrange(p32.1, p32.2, p32.3, ncol=3, nrow=1, common.legend
= TRUE, legend="bottom")
368

```